

NLP - IberLEF2026 HAHA: Humor Detection

Zuzana Neupauerová and Matej Gornál

Faculty of Mathematics, Physics and Informatics

Comenius University in Bratislava

Mlynská dolina, 842 48 Bratislava, Slovakia

{neupauerova22, gorna11}@uniba.sk

Abstract

Computational humor analysis remains a highly challenging yet rapidly evolving area within Computational Linguistics and Natural Language Processing. While significant progress has been made in automatic humor detection and classification through evaluation campaigns such as HAHA, the tasks of comprehensive humor recognition and generation are far from being fully resolved. This paper addresses the current challenges of processing figurative language by participating in the IberLEF 2026 HAHA evaluation campaign. We focus on problem of humor detection, specifically assessing the capabilities of Large Language Models (LLMs) to produce evaluating methods to automatically distinguish between human-authored and machine-generated humor. Our work aims to advance the understanding of automated humor mechanisms and contribute to the development of robust models for highly subjective NLP tasks.

1 Methodology

All code, model configurations, and replication scripts developed for our experiments are publicly available on GitHub.¹

2 Tasks

In this work, we focus on two core subtasks proposed in the IberLEF 2026 HAHA evaluation campaign, aimed at advancing the understanding of computational humor and machine-generated content:

- **Humor Detection:** determining if a news headline is a joke or not (satirical news headlines vs real news headlines). The main performance metric of this subtask will be the F1 score of the 'humorous' class. This subtask

is similar to the first subtask proposed in previous editions of the HAHA shared task, but this time it's applied to a particular domain where humorous and non-humorous content might sometimes be difficult to tell apart.

- **LLM-Generated Humor Detection:** determining if a joke inspired by a news headline was generated by an LLM or written by a human. The main performance metric for this subtask will be the F1 score of the 'automatic' class.

3 Dataset

The dataset provided for the evaluation campaign is distributed across multiple .tsv (Tab-Separated Values) files. All samples within the dataset are written in **Spanish**. Each individual sample represents a specific news headline accompanied by additional context.

The structure of the dataset consists of the following primary fields described in Table 1.

Field	Description
id	Unique identifier of the sample.
country	Country of origin where the headline was published.
date	Publication date of the headline.
headline	The main text of the news headline.
context	Additional text providing background information.

Table 1: Description of the main fields available in the dataset samples.

To monitor model performance and allow for rigorous evaluation, the data is split into exploratory (*trial*) and development (*dev*) sets for both tasks. The exact distribution of samples across the individual files is summarized in Table 2.

It is worth noting that the provided trial datasets (task1_trial.tsv and task2_trial.tsv) are exceptionally small, consisting of only 24 and 12 samples respectively. While these small subsets are highly valuable for initial pipeline validation, code

¹<https://github.com/Matejj2002/NLP-Project/tree/main>

File Name	Subtask	Number of Rows
task1_trial.tsv	Humor Detection	24
task1_dev.tsv	Humor Detection	540
task2_trial.tsv	Human AI	12
task2_dev.tsv	Human AI	350

Table 2: Data split sizes and sample distribution for Task 1 and Task 2.

debugging, and quick sanity checks of the models, their limited size poses a significant challenge for robust statistical generalization. Consequently, any performance metrics obtained on these trial sets must be interpreted with caution due to the high variance inherent to such small sample sizes. To mitigate this limitation, our primary local evaluation and hyperparameter tuning rely heavily on the larger development sets (task1_dev.tsv and task2_dev.tsv).

4 Task 1: Humor Detection

4.1 Data Analysis

We performed a brief exploratory data analysis on the task1_trial.tsv dataset to identify basic structural properties of the text.

As shown in Figure 1(a), the trial dataset is perfectly balanced, containing exactly 12 real and 12 satirical headlines. Figure 1(b) illustrates that the dataset is evenly distributed among six Spanish-speaking nations (Colombia, Mexico, Puerto Rico, Spain, Uruguay, and Venezuela), with each country providing exactly 2 real and 2 satirical samples.

Furthermore, we analyzed the text length and linguistic features across classes:

- **Text Length:** Figure 2(a) shows that satirical headlines display a wider distribution in total character length compared to real news.
- **Word Length:** As visualized in Figure 2(b), the average word length for genuine news is strictly clustered around 5.2 characters, signaling formal vocabulary. Conversely, satirical text shows much higher variance (ranging from 3.5 to over 6.0 characters per word), reflecting a more flexible and creative use of language.

4.2 Approach 1: Stable Zero-Shot Ensemble

Given the limited size of the trial dataset, fine-tuning large-scale neural networks presents a severe risk of overfitting. To mitigate this issue, we developed a training-free **Zero-Shot Classification**

Ensemble pipeline using multiple state-of-the-art cross-lingual and English models (including variants of mDeBERTa-v3, DeBERTa-v3, BART-large, and XLM-RoBERTa).


The core mechanics of this approach are designed to maximize stability and generalization across domains:

- **Prompt and Multiple Label Sets:** To capture the multi-faceted nature of humor, we evaluated each headline against three distinct label pairs mapping real vs. satirical text (e.g., “*serious journalism*” vs. “*ironic satire*”) using a standard hypothesis template: “*This text is {}.*”
- **Min-Max Score Normalization:** Since different models output raw confidence scores at varying scales, we applied *Min-Max Scaling* fitted on the trial distribution. This transformation guarantees that all models operate within a standardized range between 0 and 1.
- **Performance-Based Model Weighting:** Rather than using simple averaging, we evaluated each model independently on the trial set. Each model’s weight in the final ensemble was set proportional to its squared macro F1-score (F_1^2), heavily favoring top-performing architectures.
- **Threshold Optimization:** The final prediction is a weighted average of the normalized scores. The decision boundary was optimized using a grid search over the trial set, yielding an optimal threshold of **0.45**.


4.2.1 Experimental Results

The performance of the optimized Zero-Shot Ensemble was evaluated locally on both available datasets. On the initial pilot set (task1_trial.tsv), the ensemble achieved a well-balanced macro F1-score of **0.83**. The detailed classification report for this configuration is presented in Table 3.

When deployed on the competitive evaluation phase over the larger development set (task1_dev.tsv), the stable zero-shot ensemble successfully sustained its predictive power, achieving a final performance of **73%** accuracy on the leaderboard target. This proves that anchoring multiple zero-shot predictors with robust normalization acts as an excellent baseline for highly subjective NLP tasks without explicit training data.




(a) Class Balance




(b) Country Representation

Figure 1: Distribution of labels and geographic origins in the Task 1 trial set.



(a) Total Character Length



(b) Average Word Length

Figure 2: Analysis of text and word length properties between real and satirical classes.

Class	Precision	Recall	F1-Score	Support
real	0.83	0.83	0.83	12
satirical	0.83	0.83	0.83	12
Accuracy			0.83	24
Macro Avg	0.83	0.83	0.83	24
Weighted Avg	0.83	0.83	0.83	24

Table 3: Classification report on the trial dataset using the stable Zero-Shot Ensemble with an optimal threshold of 0.45.

4.3 Approach 2: Pseudo-Labeling and Transformer Fine-Tuning

To overcome the severe data scarcity of the trial set, we implemented a semi-supervised pipeline leveraging **Confidence-Based Pseudo-Labeling** combined with supervised transformer fine-tuning. This approach allows us to exploit the larger, unlabelled development dataset to train a robust classification model.

The pipeline operates through the following key stages:

- **High-Confidence Selection:** We first utilize

the soft probability scores generated by the stable Zero-Shot Ensemble (Approach 1) on the development set (task1_dev.tsv). To minimize the injection of noise, we apply a strict margin constraint of ± 0.15 around the optimal threshold ($t = 0.45$). Only samples falling outside this uncertainty zone (i.e., scores > 0.60 or < 0.30) are selected.

- **Dataset Augmentation:** The selected high-confidence samples are assigned their respective pseudo-labels (real or satirical) and merged with the original gold-standard trial dataset (task1_trial.tsv). This process scales up the available training data, making supervised training viable.
- **Model Architecture:** We adopt microsoft/mdeberta-v3-base as our core classifier. This model utilizes DeBERTa’s disentangled attention mechanism and is pre-trained via a multilingual masked language modeling objective, making it

exceptionally proficient at capturing fine-grained stylistic nuances like satire and humor in Spanish.

- **Training and Validation Strategy:** The augmented dataset is tokenized with a maximum sequence length of 256 tokens. The network is fine-tuned for 3 epochs with a learning rate of 2×10^{-5} , weight decay of 0.01, and a batch size of 8. Crucially, to prevent the model from overfitting to potential errors in the pseudo-labels, we use the 100% clean trial dataset as our validation set, restoring the best performing checkpoint at the end of training.
- **Final Predictions and Overfitting:** The fine-tuned model was used to execute the final inference over the development set, generating the `task1.tsv` submission file. However, empirical evaluation revealed that despite a steady loss reduction during training, this approach suffered from severe **overfitting**. Due to the propagation of pseudo-label noise and the extreme scarcity of gold-standard data, the fine-tuned model failed to generalize well, ultimately delivering lower performance on the development set compared to the training-free Zero-Shot Ensemble.